

## Task 1 :

### Virtual Memory Basics:

**Summarize how virtual memory works in Linux in three sentences or less. Why is virtual memory critical in a multitasking operating system?**

Virtual memory in Linux provides an abstraction layer that allows processes to use more memory than physically available by mapping virtual addresses to physical addresses, using both RAM and disk storage. It manages memory efficiently through paging, swapping inactive pages to disk when necessary. This enables isolation, efficient memory use, and stability in a multitasking environment.

Virtual memory is critical in a multitasking operating system because it enables multiple processes to run simultaneously without interference, optimizes the use of available physical memory, and enhances system stability and security. By providing each process with its own virtual address space, it isolates processes from one another, preventing them from accessing each other's memory, which enhances security and stability. Additionally, virtual memory allows the system to use disk space to extend RAM, ensuring that active processes have the memory they need, even if the total memory demand exceeds the physical RAM available.

## Task 2 :

### Kernel vs. User Space Memory:

**Write down the differences between kernel space and user space memory. Why is this distinction important in Linux systems?**

#### 1. Purpose and Access:

- **Kernel Space:** Reserved for the operating system kernel and its extensions. It has unrestricted access to all system resources and hardware.
- **User Space:** Used by user applications and processes. It has restricted access to system resources and cannot directly access hardware.

#### 2. Memory Protection:

- **Kernel Space:** Runs in a privileged mode (ring 0) with full access to hardware and memory. Errors in kernel space can crash the entire system.
- **User Space:** Runs in a non-privileged mode (ring 3) with limited access. Errors in user space typically only affect the offending process.

#### 3. Memory Addressing:

- **Kernel Space:** Part of the virtual memory space allocated to the kernel. It is mapped to high memory addresses.

- **User Space:** Part of the virtual memory space allocated to user applications. It is mapped to low memory addresses.

#### 4. Performance:

- **Kernel Space:** Direct access to hardware and resources makes operations faster but requires careful handling to avoid system instability.
- **User Space:** Indirect access to hardware and resources through system calls, which adds overhead but provides safety and stability.

#### 5. System Calls:

- **Kernel Space:** Does not require system calls to execute privileged operations as it operates with elevated privileges.
- **User Space:** Must use system calls to request kernel services, which switch execution from user mode to kernel mode.

#### 1. Stability:

- **Kernel Space:** Errors in the kernel space can affect the entire system because the kernel has unrestricted access to all hardware and memory. Keeping user applications out of kernel space helps to protect the system from crashes caused by user-level errors.
- **User Space:** Errors in user space generally only affect the specific application that caused them, leaving the rest of the system running smoothly.

#### 2. Security:

- **Kernel Space:** The kernel has full control over system resources and hardware. Restricting access to kernel space prevents malicious or buggy code from compromising the entire system.
- **User Space:** User space is isolated from kernel space, meaning user applications cannot directly access critical system resources. This isolation helps protect the system from attacks and accidental damage.

#### 3. Performance:

- **Kernel Space:** Direct access to hardware allows the kernel to perform critical tasks efficiently. However, keeping most applications in user space reduces the risk of performance issues caused by poorly written code.
- **User Space:** User applications can perform their tasks without risking the stability of the kernel, although they might incur a slight performance overhead due to the need for system calls to access hardware or system resources.

#### 4. Resource Management:

- **Kernel Space:** The kernel manages system resources such as CPU, memory, and I/O devices. By controlling access to these resources, the kernel can ensure fair distribution and prevent any single process from monopolizing them.
- **User Space:** User processes request resources from the kernel, which manages and allocates these resources efficiently, preventing conflicts and ensuring that each process gets the resources it needs.

## 5. Development and Debugging:

- **Kernel Space:** Developing and debugging kernel code is complex and risky since errors can crash the entire system. The kernel space is reserved for core system functions, which need to be robust and reliable.
- **User Space:** User space is more forgiving for development and debugging. Issues in user applications can be resolved without affecting the overall system, making it a safer environment for developers.

## Task 3 :

### Exploring procfs:

**What kind of information can you find about memory mappings in the procfs?  
Describe how you would use procfs to find out about the virtual memory**

The /proc filesystem (procfs) in Linux provides a wealth of information about memory mappings for processes. Specifically, you can find detailed memory mapping information in the /proc/[pid]/maps and /proc/[pid]/smaps files, where [pid] is the process ID of the process you are interested in. Here's what you can find in these files:

#### 1. /proc/[pid]/maps:

- **Address Ranges:** The virtual memory address range of each memory region.
- **Permissions:** Access permissions of the memory region (read, write, execute, shared/private).
- **Offset:** Offset into the file or the device if the memory region is mapped to a file.
- **Device:** Device number in hexadecimal format if the memory region is mapped to a file.
- **Inode:** Inode number of the mapped file if applicable.
- **Pathname:** Path to the mapped file or region (e.g., shared libraries, heap, stack).

#### 2. /proc/[pid]/smaps:

- **Detailed Memory Usage:** Provides more granular information about each memory region, including:
  - **Size:** Total size of the memory region.
  - **Rss:** Resident set size, the portion of memory currently in RAM.
  - **Pss:** Proportional set size, the portion of memory shared with other processes.
  - **Shared\_Clean:** Clean shared memory.
  - **Shared\_Dirty:** Dirty shared memory.
  - **Private\_Clean:** Clean private memory.
  - **Private\_Dirty:** Dirty private memory.
  - **Referenced:** Amount of memory that has been referenced.

- **Anonymous:** Amount of anonymous memory.
- **Swap:** Amount of memory swapped out to disk.
- **KernelPageSize and MMUPageSize:** Page sizes.
- **Locked:** Whether the memory region is locked in RAM.

#### 1. Identify the Process ID (PID):

- First, determine the PID of the process you are interested in. You can use commands like `ps`, `top`, or `pidof` to find this information. For example, `ps aux | grep [process_name]` can help you find the PID of a process by name.

#### 2. Check the Memory Mappings:

- Navigate to the `/proc` directory for the specific process using its PID. For example, if the PID is 1234, you will look in `/proc/1234`.

#### 3. Examine `/proc/[pid]/maps`:

- View the memory mappings with the following command

```
cat /proc/1234/maps
```

- This file shows the address ranges, permissions, offset, device, inode, and pathname of the memory regions. It provides a basic overview of how the process's virtual memory is laid out.

#### Examine `/proc/[pid]/smaps`:

- For more detailed information, view the `smaps` file

```
cat /proc/1234/smaps
```

- This file gives comprehensive details about each memory region, including size, resident set size (RSS), proportional set size (PSS), shared/private clean and dirty memory, referenced memory, anonymous memory, swap usage, page sizes, and whether the memory is locked.

#### 4 . Check Memory Usage Summary:

- You can also look at `/proc/[pid]/status` to get a summary of memory usage

```
cat /proc/1234/status
```

## Task 4 :

### Buffer Overruns Prevention:

Explain what buffer overruns are and how the Linux kernel prevents them. Why is it important to avoid buffer overruns in relation to virtual memory?

### Buffer Overruns Explained

Buffer overrun, also known as buffer overflow, occurs when a program writes more data to a buffer than it can hold. This extra data can overwrite adjacent memory, leading to unpredictable behavior, crashes, or security vulnerabilities. Buffer overruns are especially dangerous in languages like C and C++ that do not perform automatic bounds checking on arrays.

### How the Linux Kernel Prevents Buffer Overruns

#### 1. Memory Protection Mechanisms:

- **Segmentation and Paging:** The Linux kernel uses hardware memory protection features, such as segmentation and paging, to isolate different memory regions. This ensures that a process cannot access memory outside its allocated segments.

#### 2. Address Space Layout Randomization (ASLR):

- **Randomizing Memory Layout:** ASLR randomizes the memory addresses used by system and application processes, making it difficult for an attacker to predict the location of specific buffers, thereby reducing the likelihood of successful exploitation of buffer overruns.

#### 3. Stack Canaries:

- **Stack-Based Protection:** The kernel uses stack canaries (special values placed between a buffer and control data on the stack) to detect buffer overruns. If a buffer overrun occurs, the canary value is altered, and the program can detect the corruption and terminate before any malicious code can be executed.
4. **Non-Executable Memory Regions:**
    - **Executable Space Protection (NX):** The kernel marks certain regions of memory (such as the stack and heap) as non-executable, preventing code from being run from these regions. This limits the ability of attackers to execute arbitrary code via buffer overrun vulnerabilities.
  5. **Compiler-Based Protections:**
    - **Fortify Source:** When compiling applications, the Linux kernel and many user-space programs use compiler flags like `-D_FORTIFY_SOURCE=2`, which adds checks to detect buffer overflows in certain functions (e.g., `strcpy`, `sprintf`).
    - **Position-Independent Executables (PIE):** PIE makes use of ASLR more effective by ensuring that not only shared libraries but also the main executable is randomized.

## Importance of Avoiding Buffer Overruns in Relation to Virtual Memory

1. **System Stability:**
  - **Prevent Crashes:** Buffer overruns can corrupt memory, leading to application crashes and system instability. In a multitasking environment, this can affect multiple processes and degrade overall system performance.
2. **Security:**
  - **Prevent Exploitation:** Buffer overruns are a common attack vector for executing arbitrary code, escalating privileges, and compromising system security. By preventing buffer overruns, the integrity and confidentiality of the system and user data are protected.
3. **Memory Integrity:**
  - **Protect Data:** Virtual memory relies on proper management of memory regions. Buffer overruns can corrupt not only the memory of the vulnerable application but also adjacent memory areas, potentially affecting other applications and the kernel itself.
4. **Isolation of Processes:**
  - **Maintain Isolation:** Virtual memory ensures that each process operates in its own memory space. Buffer overruns can break this isolation by allowing one process to inadvertently or maliciously modify the memory of another process.

